

Fundamentals of Artificial Intelligence

Lecture-2 Problem Solving

Debela Desalegn


March 16, 2025

Today we will see:

- Search Problems
- Uninformed Search Methods
 - Depth-First Search (DFS)
 - ✓ Breadth-First Search (BFS)
 - ✓ Uniform-Cost Search(UCS)

Intro

- **Problem Solving Agent:** An intelligent agent that finds solutions by searching for the best sequence of actions.
- **Search Algorithms:** Informed vs Uninformed
- **Assumption:** Known world
- **Problem Solving Phases**
 - Goal Formulation
 - Problem Formulation (Abstraction)
 - Search (simulated)
 - Execution (Closed Loop vs Open Loop)




Intro

Reflex agents:

- Choose action based on current percept (and maybe memory)
- May have memory or a model of the world's current state
- Do not consider the future consequences of their actions
- **Consider how the world is ?**

Can a reflex agent be rational?




Intro

Planning agents:

- Ask “what if”
- Decisions based on (hypothesized) consequences of actions
- Must have a model of how the world evolves in response to actions
- Must formulate a goal (test)
- **Consider how the world WOULD BE?**

Optimal vs. complete planning

Planning vs. replanning



Search Problems

State: Represents a configuration of the problem.

- **Initial State:** Where the agent starts.
- **Goal State:** The desired outcome.
- **State Space:** The set of all possible states reachable from the initial state.
- **Actions:** The possible moves the agent can take (e.g., moving a piece in a puzzle, turning left in a navigation problem).
- **Transition Model:** Defines how an action transforms the current state into a new state.
- **ACTION-COST**(s, a, s') or $c(s, a, s')$: Measures the cost of performing action a in state s to reach state s' .
- **Path:** A sequence of actions leading from the initial state to the other state.
- **Solution:** a sequence of actions (a plan/path) which transforms the start state to a goal state
- **Optimal Solution == Lowest ACTION-COST**


Search Problems

A search problem consists of:

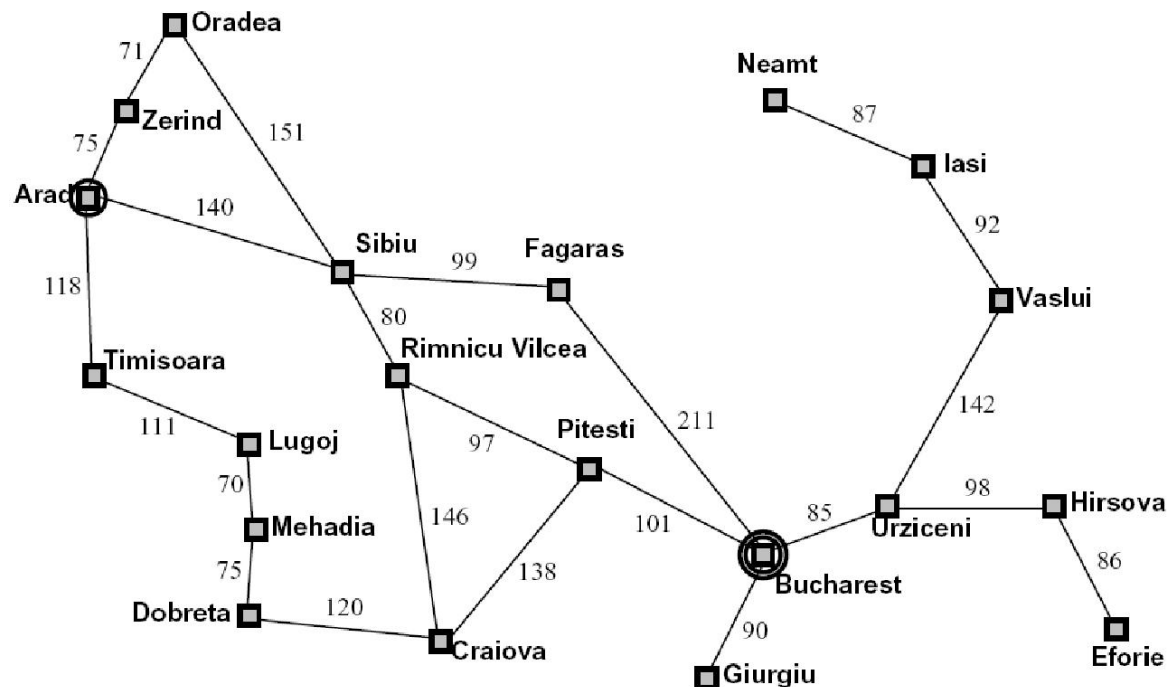
✓ A state space



✓ A successor function (with actions, costs)



✓ A start state and a goal test



State space:

- Cities

Successor function:

- Roads: Go to adjacent
- city with cost = distance

Start state:

- Arad

Goal test:

- Is state == Bucharest?

Solution?

What's in a State Space?

The world state includes every last detail of the environment.

A search state/search tree keeps only the details needed for planning (abstraction).

Problem: Pathing

- States: (x,y) location
- Actions: NSEW
- Successor: update location only
- Goal test: is (x,y)=END


Problem: Eat-All-Dots

- States: {(x,y), dot booleans}
- Actions: NSEW
- Successor: update location and possibly a dot boolean
- Goal test: dots all false

State Space Graphs

State space graph: A mathematical representation of a search problem

- **Nodes** are (abstracted) world configurations
 - **Arcs** represent successors (action results)
 - The **goal test** is a set of goal nodes (maybe only one)
- ▮ In a state space graph, each state occurs only once!
- We can rarely build this full graph in memory (it's too big), but it's a useful idea




State Space Graphs

State space graph: A mathematical representation of a search problem

- **Nodes** are (abstracted) world configurations
- **Arcs** represent successors (action results)
- The **goal test** is a set of goal nodes (maybe only one)


▮ In a state space graph, each state occurs only once!

- We can rarely build this full graph in memory (it's too big), but it's a useful idea



Tiny state space graph for a tiny search problem

Search Trees




A search tree:

- A “what if” tree of plans and their outcomes
- The start state is the root node
- Children correspond to successors
- Nodes show states, but correspond to PLANS that achieve those states
- For most problems, we can never actually build the whole tree

State Space Graphs vs. Search Trees


State Space Graph



Each NODE in the search tree is an entire PATH in the state space graph.


We construct both on demand – and we construct as little as possible.

Search Tree

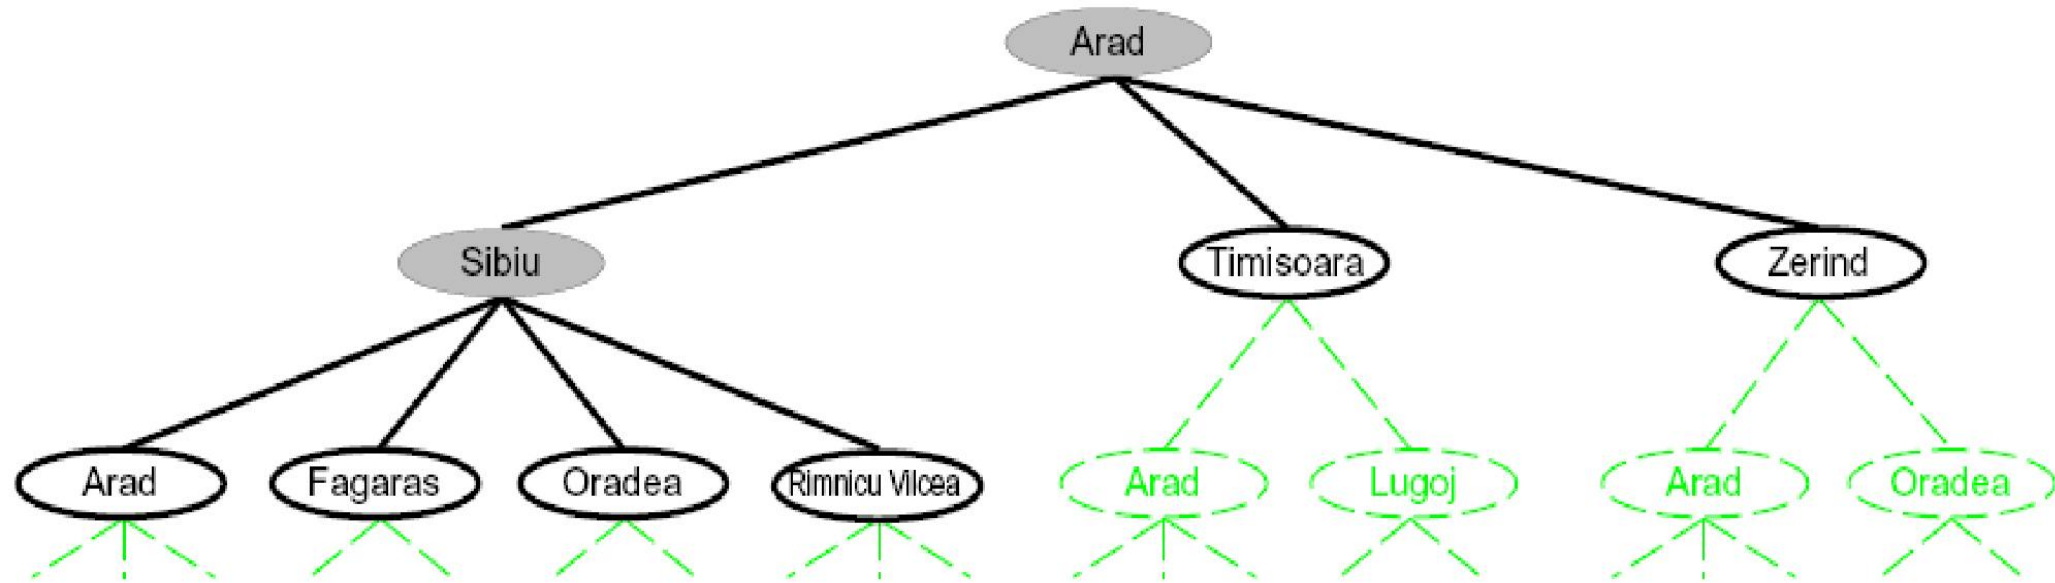


Quiz-1

How big is its search tree (from S)?



Tree Search



Search

- Expand out potential plans (tree nodes)
- Maintain a **fringe/frontier** of partial plans under consideration
- Try to expand as few tree nodes as possible

General Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```



Important ideas:

- Fringe/Frontier
- Expansion
- Exploration strategy

Main question: Which fringe nodes to explore?

The essence of search

Example: Tree Search



~~s~~
~~s → d~~
 s → e
 s → p
 s → d → b
 s → d → c
~~s → d → e~~
 s → d → e → h
~~s → d → e → r~~
~~s → d → e → r → f~~
 s → d → e → r → f → c
~~s → d → e → r → f → G~~

Search Data structure

- **node.STATE**: the state to which the node corresponds;
- **node.PARENT**: the node in the tree that generated this node;
- **node.ACTION**: the action that was applied to the parent's state to generate this node;
- **node.PATH-COST**: the total cost of the path from the initial state to this node. In mathematical formulas, we use $g(n)$ as a synonym for PATH-COST.
- A data structure for a frontier could be a Queue (Priority Queue, FIFO Queue, LIFO queue (stack)).

Uninformed Search Strategies

No clue about how close a state is to the goal(s).

Have information on how to traverse or visit the nodes in the tree.

- ✓ Depth-First Search (DFS)
- ✓ Breadth-First Search (BFS)
- ✓ Uniform-Cost Search(UCS)



Depth-First Search (DFS)

Strategy: expand a deepest node first

Implementation: Fringe is a LIFO stack



Depth-First Search (DFS)



Depth-First Search (DFS) Properties

What nodes DFS expand?

- Some left prefix of the tree.
- Could process the whole tree!
- If m is finite, takes time $O(b^m)$

How much space does the fringe take?


- Only has siblings on path to root, so $O(bm)$

Is it complete?

- m could be infinite, so only if we prevent cycles (more later)
- Incomplete: in infinite state spaces



Is it optimal?

- No, it finds the “leftmost” solution, regardless of depth or cost



Breadth-First Search (BFS)

- Optimal when all action cost is the same.
- Node Expansion sequence
 - Root, successors of the root, next successors
 - Complete in infinite state spaces.
 - Is Best first search where $f(n)$ is the depth.
 - Good efficiency \rightarrow FIFO queue





Breadth-first search on a simple binary tree. At each stage, the node to be expanded next is indicated by the triangular marker.

Breadth-First Search (BFS)

Strategy: expand a shallowest node first

Implementation: Fringe is a FIFO queue



Breadth-First Search (BFS)

```
function BREADTH-FIRST-SEARCH(problem) returns a solution node or failure  
  node  $\leftarrow$  NODE(problem.INITIAL)  
  if problem.IS-GOAL(node.STATE) then return node  
  frontier  $\leftarrow$  a FIFO queue, with node as an element  
  reached  $\leftarrow$  {problem.INITIAL}  
  while not IS-EMPTY(frontier) do  
    node  $\leftarrow$  POP(frontier)  
    for each child in EXPAND(problem, node) do  
      s  $\leftarrow$  child.STATE  
      if problem.IS-GOAL(s) then return child  
      if s is not in reached then  
        add s to reached  
        add child to frontier  
  return failure
```


Breadth-First Search (BFS)

What nodes does BFS expand?

- Processes all nodes above shallowest solution
- Let depth of shallowest solution be s
- Search takes time $O(b^s)$

How much space does the fringe take?


- Has roughly the last tier, so $O(b^s)$
- The memory requirements are higher than time requirements.

Is it complete?

- s must be finite if a solution exists, so yes!


Is it optimal?

- Only if costs are all 1 (more on costs later)



Quiz 2: DFS vs BFS

- | When will BFS outperform DFS?
- | When will DFS outperform BFS?



Iterative Deepening

Depth Limited Depth-First Search


- Limit depth to some specific level “l” and get a time and space complexity of $O(b^l)$ and $O(bl)$ respectively.

Idea: get DFS’s space advantage with BFS’s time / shallow-solution advantages

- Run a DFS with depth limit 1. If no solution...
- Run a DFS with depth limit 2. If no solution...
- Run a DFS with depth limit 3.

Isn’t that wastefully redundant?

- Generally most work happens in the lowest level searched, so not so bad!
- Good if the diameter of the problem is known



Iterative Deepening

Time Complexity: $O(b^d)$ when there is a solution and $O(b^M)$ when there is none.

Space Complexity: $O(bd)$ when there is a solution and $O(bm)$ when there is none.

function ITERATIVE-DEEPENING-SEARCH(*problem*) **returns** a solution node or *failure*

for *depth* = 0 **to** ∞ **do**

result \leftarrow DEPTH-LIMITED-SEARCH(*problem*, *depth*)

if *result* \neq *cutoff* **then return** *result*

function DEPTH-LIMITED-SEARCH(*problem*, ℓ) **returns** a node or *failure* or *cutoff*

frontier \leftarrow a LIFO queue (stack) with NODE(*problem*.INITIAL) as an element

result \leftarrow *failure*

while not IS-EMPTY(*frontier*) **do**

node \leftarrow POP(*frontier*)

if *problem*.IS-GOAL(*node*.STATE) **then return** *node*

if DEPTH(*node*) > ℓ **then**

result \leftarrow *cutoff*

else if not IS-CYCLE(*node*) **do**

for each *child* **in** EXPAND(*problem*, *node*) **do**

 add *child* to *frontier*

return *result*

Iterative Deepening

Goal to M





limit: 0















limit: 1



limit: 2





limit: 3



Dijkstra's algorithm or Uniform Cost Search

Strategy: expand a cheapest node first:

Fringe is a priority queue (priority: cumulative cost)



Dijkstra's algorithm or Uniform Cost Search

What nodes does UCS expand?

- Processes all nodes with cost less than cheapest solution!
- If that solution costs C^* and arcs cost at least ε , then the “effective depth” is roughly C^*/ε
- Takes time $O(b^{C^*/\varepsilon})$ (exponential in effective depth)

How much space does the fringe take?


- Has roughly the last tier, so $O(b^{C^*/\varepsilon})$

Is it complete?

- Assuming best solution has a finite cost and minimum arc cost is positive, yes!


Is it optimal?

- Yes! (Proof next lecture via A^*)



Example

Find the shortest path from state A to state F.



Next Lecture: Searching Problem

Informed Search

Heuristics (Greedy Search and A* Search)

Graph Search